


# Intro to Deep Learning and Computer Vision

MIT GSL-PRO, Uruguay 2020  
Week 3, Day 1

# How Does Deep Learning Relate to Previous Models We've Covered?


- Let's revisit our **logistic regression** model:



- We can think of this as a **single unit neural network** with a **sigmoid activation**!


# Neural Networks As Cascaded Groupings of Units

- When **F** is a **sigmoid**, this becomes logistic regression.
- When **F** is **linear**, this becomes linear regression.
- When the **F** is **cross-entropy**, this becomes binary/multiclass classification.




# Neural Networks As Cascaded Groupings of Units

- This intuition will be helpful as we discuss how networks **make predictions**, and how they are **trained**.



# Review from Last Week – Neural Networks

- Last week, we discussed **neural networks**, and how they can be used for **supervised learning**.
- In addition to **NLP** and **e-commerce**, neural networks can also be used to solve problems in **computer vision** and **healthcare**.



# Applications of Neural Networks

## Autonomous Vehicles



## Google



## Insurance Claims



## Chatbots




# Different Classes of Neural Networks

Many different kinds of neural networks:


- **Feedforward**
- **Convolutional (CNNs)**
- **Recurrent (RNNs)**
- **Graph**
- **Variational Autoencoders\* (VAEs)**
- **Generative Adversarial Networks (GANs)**

This course only analyzes the first two kinds of networks.


[1] Feedforward network class



[2] Convolutional network class



[3] Recurrent network class




\* VAEs are actually unsupervised! They are part of a class of algorithms known as **deep generative models**.

# Anatomy of a Feedforward Neural Network


Feedforward neural networks have 3 basic elements:

- **Neurons**
- **Weights**
- **Activations**



Feedforward networks are composed of **cascaded layers** of neurons, weights, and activations


- **Input layer**
- **Hidden layer(s)**
- **Output layer**





# How Neural Networks Make Predictions

- Input is propagated through the network with the **forward** algorithm
- At each layer:
  - Neuron **activations** from previous layer are multiplied by layer **weights**, and **summed** together.
  - Activation function is applied to the **sum**.



[1] Weighted sum of activations:

$$z^{(L)}_i = \sum_j x_j w^{(L-1)}_{ji} + w^{(L-1)}_{0i}$$


[2] Activation of sum:

$$a^{(L)}_i = f(z^{(L)}_i)$$


# Intuition for Neural Network Predictions

- Trained neural networks learn “when activations should be **high** or **low**”
- Output layers typically reflect what the network is used for:
  - Any real number output → **Regression**
  - $\{1, 0\}$  output → **Binary Classification**
  - $\{1, 2, 3, \dots, 10\}$  output → **Multiclass Classification**


Binary classification



MULTI-class classification




Linear Regression with Neural Networks!



# Output Activation Functions and their Applications


## [1] Regression:

- Linear:  $f(x) = x$ , varies from  $(-\infty, \infty)$
- Rectified Linear Unit (ReLU):  $\text{ReLU}(x)$ , varies from  $(0, \infty)$



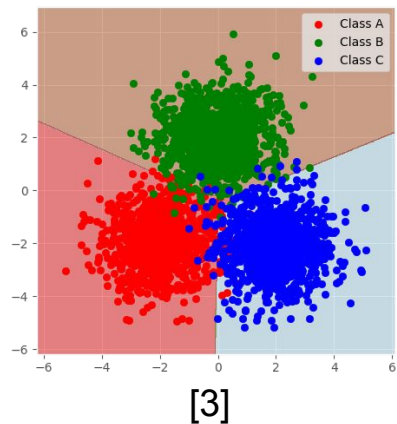
## [2] Logistic Regression/Binary Classification:

- Sigmoid:  $\sigma(x)$ , varies from  $(0,1)$
- Hyperbolic Tangent:  $\tanh(x)$ , varies from  $(-1,1)$



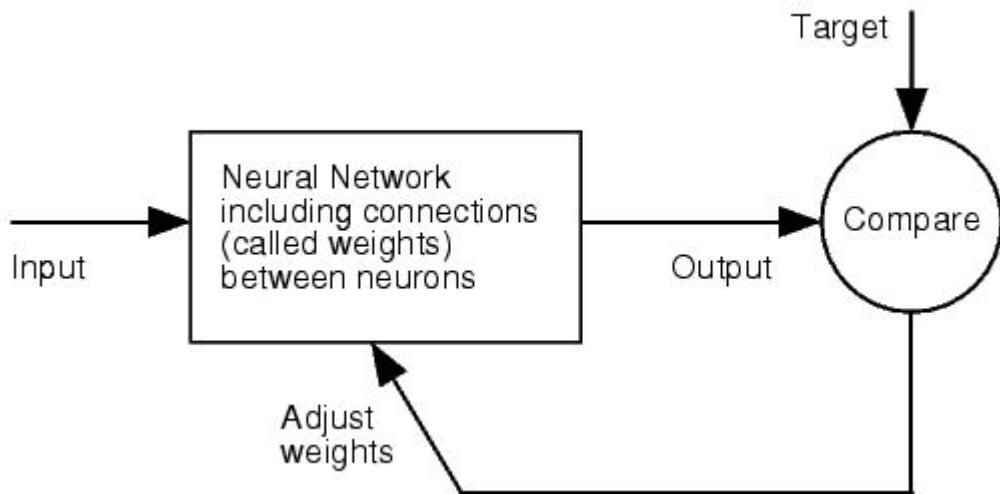
## [3] Multiclass Classification

- SoftMax:  $\text{SM}(x_i)$ , varies from  $(0,1)$



# Training Neural Networks – Overview



- For neural networks to be useful, we need to **train** them
- Neural networks can be trained by **optimizing** their **weights**.
  - This is how neural networks “**learn**”!
  - Different **weights** lead to different **predictions**
- These weights are “learned” through the **backpropagation** algorithm.



# The Backpropagation Algorithm

Algorithm composed of two stages:





1. **Forward pass:** Network makes prediction with labeled training data.
2. **Backward pass:** Network weights are updated according to how “wrong” the prediction is.



# Overview – Deep Learning

Field of **Deep Learning** uses **Deep Neural Networks (DNNs)**, which can learn remarkably complicated tasks, given sufficient data and training. Some examples include:

- Object **detection** and **classification**
- Speech and image **generation**
- High-dimensional **predictions**
- **Recommendation** systems
- **Decision-making** for agents




# Deep Learning Is A Very New Field

Field has seen a lot of advances in recent years from development of better **CPUs** and graphical/tensor processing units (**GPUs/TPUs**).



Advances are quite literally happening every day!

Deep Learning Publications by Year



# Deeper, Not Wider, Neural Networks\*

more layers with less weights > less layers with more weights




\*For More Information on this: [Stack Overflow Post](#).



# Why Deep Learning?

- Why should we even consider deep learning when we already have machine learning models that can solve our problems?
- Deep learning can often achieve **higher performance** than other models
- Caveat: **Need more data** in order to achieve **higher performance**



# Deep Learning Packages in Python

These packages allow us to implement, train, and evaluate million-parameter neural networks in **less than 15 lines of code!**

Main deep learning package we will use in this class: **Keras**.

Other packages to explore:

- TensorFlow
- PyTorch
- Caffe
- Chainer
- FastAI



Keras



Chainer

PYTORCH

fast.ai

Caffe  
MODELS



TensorFlow

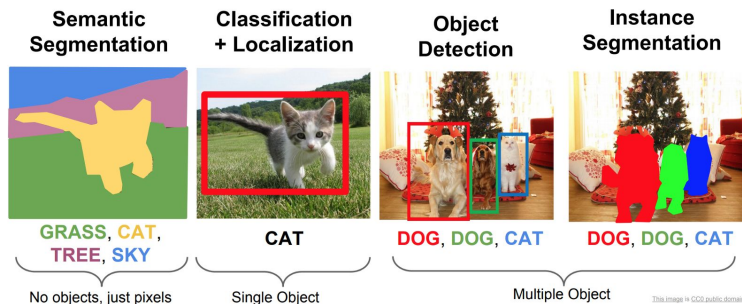
# **Week 3, Day 2**

# **Intro to Computer Vision**

MIT GSL-PRO, Uruguay 2020


# Overview – Computer Vision

- **Computer Vision** is a subset of artificial intelligence primarily concerned with understanding **spatial and imagery** data. Typical applications include:
  - Image blurring, sharpening, and edge detection
  - Capturing motion
  - Image classification/object detection and classification/scene segmentation
  - Image and video generation
- This field developed independently of deep neural networks, but has improved substantially over the last decade through integration of these deep network frameworks.



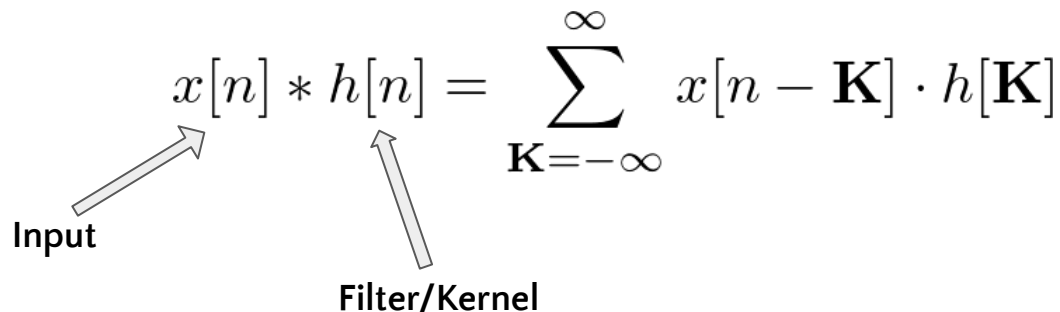
# Computer Vision as Filtering

- **Filtering** is the most central idea in **computer vision**.
- Intuitive idea behind filtering: A **system F** transforms an **input** into an **output**.



# Filtering – Convolutions

- Filtering is done through **convolution** and **kernels**
- **Convolution** is a mathematical operation that computes a **sum of linear combinations of the input**


$$x[n] * h[n] = \sum_{\mathbf{K}=-\infty}^{\infty} x[n - \mathbf{K}] \cdot h[\mathbf{K}]$$


The diagram illustrates the convolution equation. An arrow labeled "Input" points to the term  $x[n]$  in the equation. Another arrow labeled "Filter/Kernel" points to the term  $h[n]$  in the equation. The equation itself is  $x[n] * h[n] = \sum_{\mathbf{K}=-\infty}^{\infty} x[n - \mathbf{K}] \cdot h[\mathbf{K}]$ .

- **Kernels** are matrices\* which we **convolve** our **input** with to **filter**!


\* For discrete problems, kernels can be matrices. For continuous problems, they take the more general form of an inner product of a function applied to the two inputs. Link [HERE](#).

# Convolution\* – Building Intuition

- Think of convolution as **sliding a filter** across the **input**, and computing the product between the **filter/kernel** and **input** for each input value
- This concept is illustrated on the next slides




$x[n]$  (Input)



$h[n]$  (Filter)

\* Technically, this is **correlation**, not **convolution**. They are both closely related, with just a simple change of sign. If interested, see the link [HERE](#).

# Convolution – 1D Example



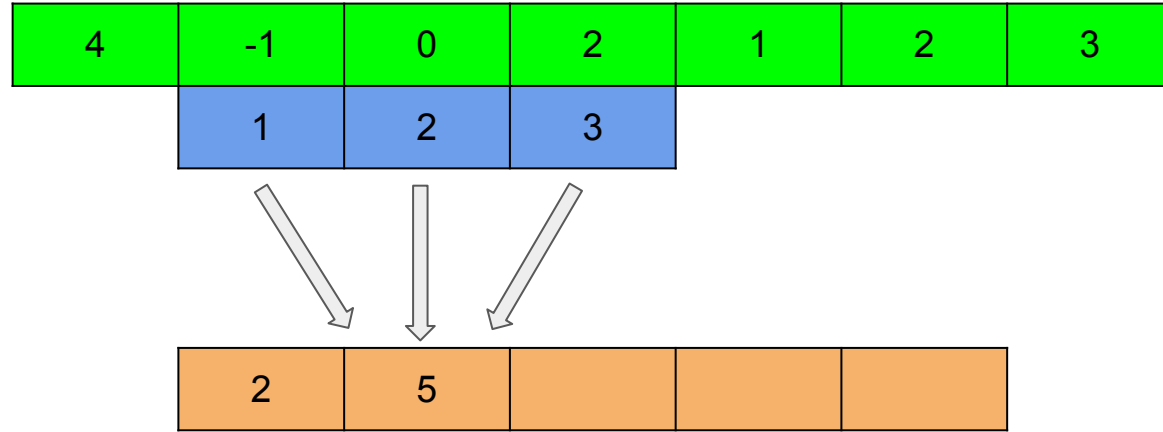
$x[n]$  (Input)

$h[n]$  (Filter)

$$y[n] = x[n] * h[-n]$$



# Convolution – 1D Example

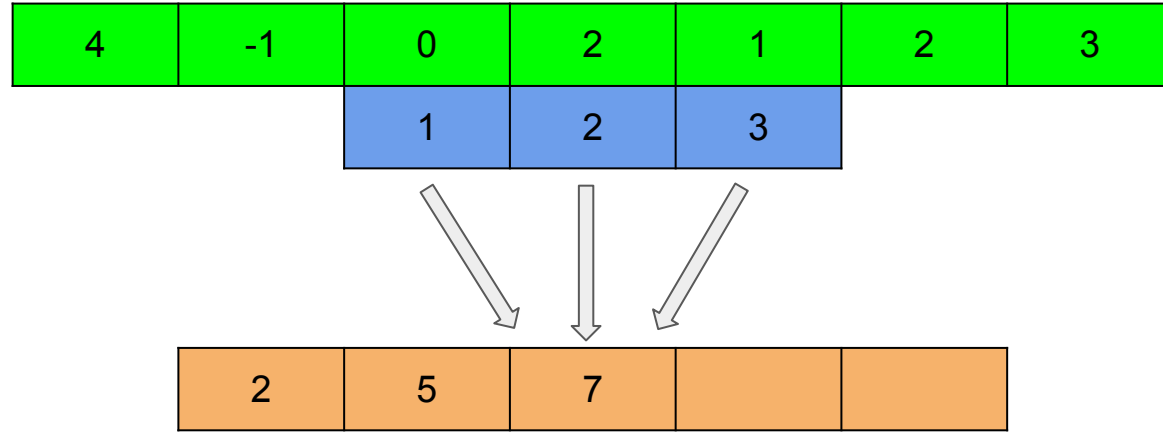


$x[n]$  (Input)

$h[n]$  (Filter)

$$y[n] = x[n] * h[-n]$$

# Convolution – 1D Example

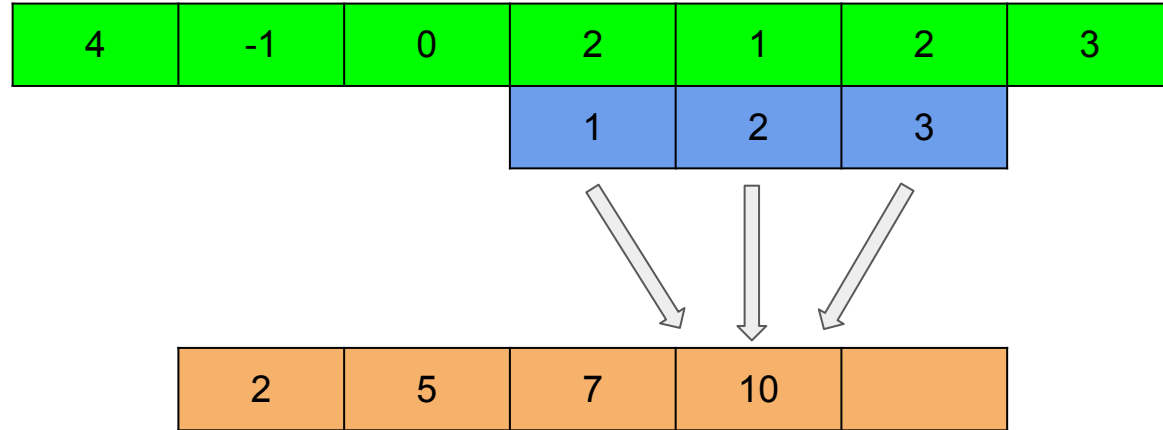


$x[n]$  (Input)

$h[n]$  (Filter)

$$y[n] = x[n] * h[-n]$$


# Convolution – 1D Example



$x[n]$  (Input)

$h[n]$  (Filter)

$$y[n] = x[n] * h[-n]$$



# Filtering – “Padding” Techniques

- What do we do if we want the output to be the same size as input?
  - We can use **padding**!
- Different Types of Padding:
  - **Valid padding** – Do nothing (does not fix size problem)


4	-1	0	2	1	2	3
---	----	---	---	---	---	---

- **Zero padding** – Add 0's to the edges of the input

0	4	-1	0	2	1	2	3	0
---	---	----	---	---	---	---	---	---


# Convolution - 2D Example

- In 2D, we use **matrices** for our **kernels**




# Convolution - 2D Example

- In 2D, we use **matrices** for our **kernels**



# Convolution - 2D Example


- In 2D, we use **matrices** for our **kernels**






# Convolution – 2D Example

- In 2D, we use **matrices** for our **kernels**




# 2D Convolution with Padding

- In 2D, we use **matrices** for our **kernels**
- **Padding with zeros** ensures the output is the same size as the input
- Notice how we do the same sum of products between elements of the kernel and input as we saw before!



- Filtering primarily used in the image domain
- Can use filters for a variety of tasks, such as:
  - **Blurring**
  - **Sharpening**
  - **Edge detection**
  - **Corner Detection**



# Filtering – Types of Kernels

- Blurring


$\frac{1}{9}$	1	1	1
	1	1	1
	1	1	1

- Gaussian Blurring


$\frac{1}{209}$	16	26	16
	26	41	26
	16	26	16



input



output




# Filtering – Types of Kernels

- Sharpening

-1	0	1
-2	0	2
-1	0	1



input



output




- Identity

0	0	0
0	1	0
0	0	0



input




output



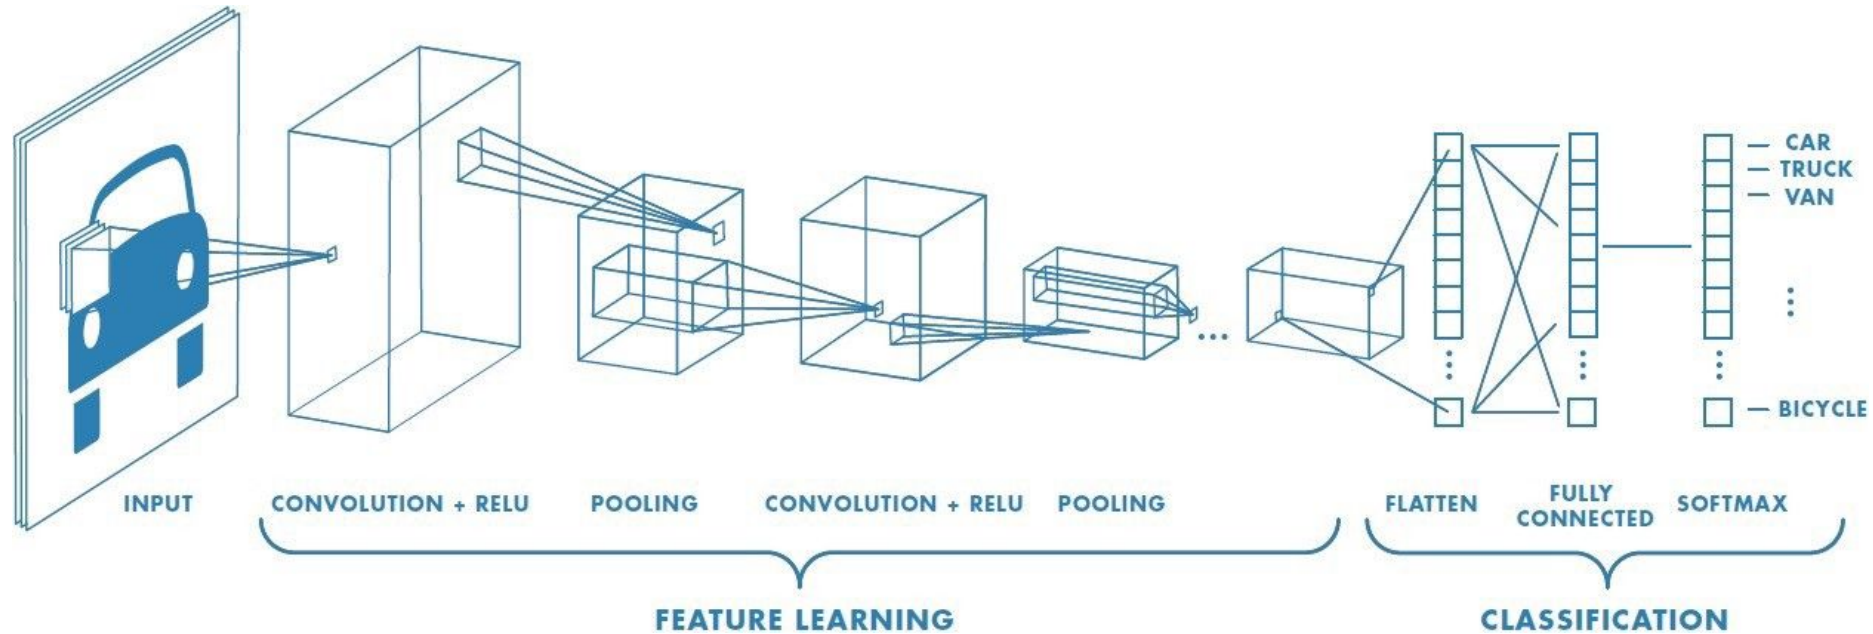
# Computer Vision with Deep Learning – CNNs

- Computer Vision connected to Deep Learning primarily through **Convolutional Neural Networks (CNNs)**.
- These networks use **convolutional layers**
  - Finds **local features** in an image
  - But can find features **anywhere** in an image (**translation invariance**)
- Convolutional layers use the same **convolution\*** operation that we saw for filtering!



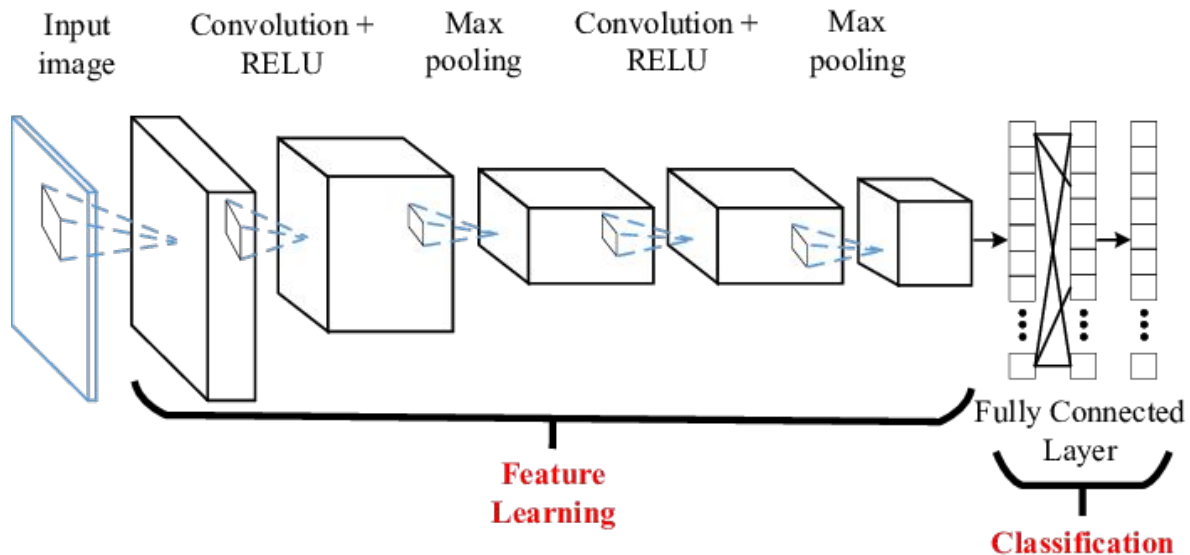
\* Technically, this is **correlation**, not **convolution**. They are both closely related, with just a simple change of sign. If interested, see the link [HERE](#).

# Anatomy of a Convolutional Neural Network (CNN)



# Core Elements of CNNs

- Input “layer”
- Convolutional layers
- Max pooling layers
- Flattening layer\*
- Fully Connected layers\*
- Output layers




\*Some neural networks, such as Fully Convolutional Neural Networks, don't use these elements.



# Introducing Non-Linearities in CNNs: Max Pooling

- Recall with feedforward networks that we primarily used **ReLU** as a means to introduce non-linearities into our neural networks.
- In CNNs, we will introduce non-linearity as well – through **max pooling** functions.



# Tasks in CV + DL

- Recent advances in automated tasks have been achieved through recent advances in CV and DL:

- Object detection
- Semantic segmentation
- Image classification
- Image and video generation




Image Classification



Object Detection



Semantic Segmentation



# Applications: Using CNNs for Vision in Healthcare

## [1] Pneumonia Prediction:

This week, we'll be using Convolutional Neural Networks to predict whether a patient has pneumonia using chest x-ray imagery.

## [2] Breast Cancer Prediction:

A hybrid Deep Learning and Logistic Regression model yielded substantial improvement in breast cancer diagnosis over only the Logistic Regression model. Paper [HERE](#).

## [3] Automated Detection of Diabetic Retinopathy and Diabetic Macular Edema:

Uses CNNs; named by JAMA as one of the best papers of the decade. Paper [HERE](#).


# ADDITIONAL SLIDES

(for students' reference)

# Variational Autoencoders (VAEs) – NNs for Unsupervised Learning

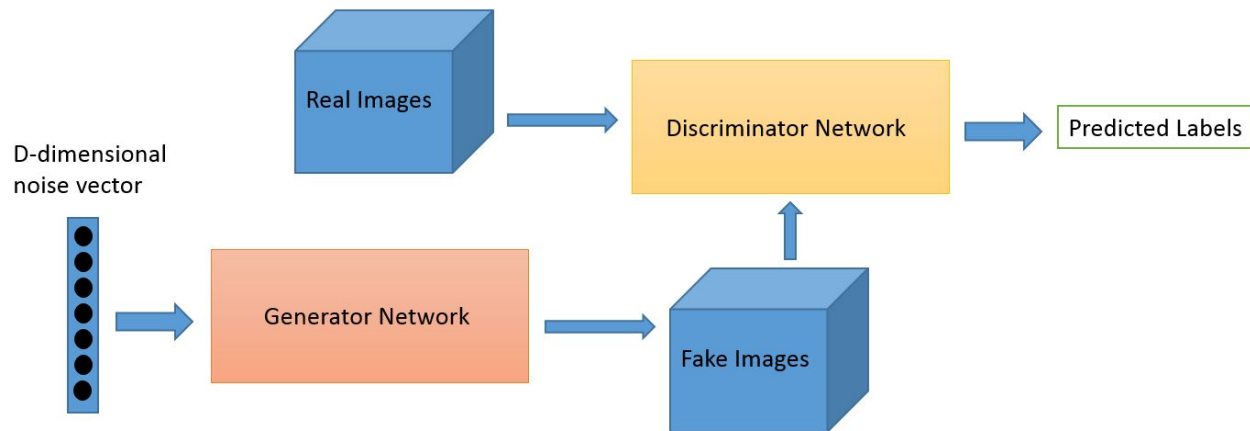
- Variants of neural networks that don't require labeled data:

## Variational Autoencoders (VAEs)



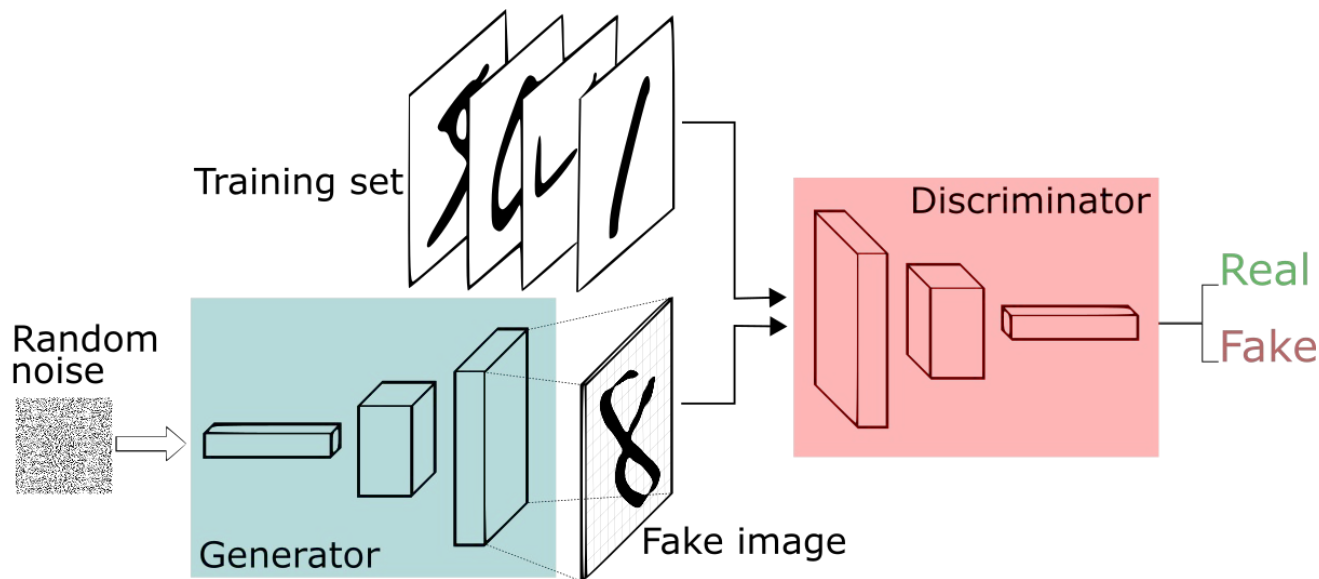
# Generative Adversarial Networks (GANs)

- Composed of **generator** and **discriminator** networks.
- **Generator** is given random noise, and tries to output fake images to fool the discriminator.
- **Discriminator** is trained on real and fake images, and tries to learn how to spot fake images.
- Practical note: Proper training of these networks is quite difficult, and requires careful tuning.



# GANs + CV: Deep Convolutional GANs (DCGAN)

- Useful link for tutorial here: [TensorFlow](https://www.tensorflow.org/tutorials/gan)
- This architecture can be used for real-world image generation!



# Classical CV Object Detection: Scale-Invariant Feature Transform (SIFT)

- CV technique for object detection that doesn't use deep learning
- Original paper can be found [HERE](#), code can be found [HERE](#)






Image gradients



Keypoint descriptor


## SIFT (Scale-Invariant Feature Transform)

- [1] Find extremes
- [2] Detect keypoints
- [3] Assign orientations
- [4] Assign keypoint descriptors

Images courtesy of David G. Lowe, "Distinctive Image Features from Scale-Invariant Keypoints."



# Template Matching: When Invariance Is Not Important

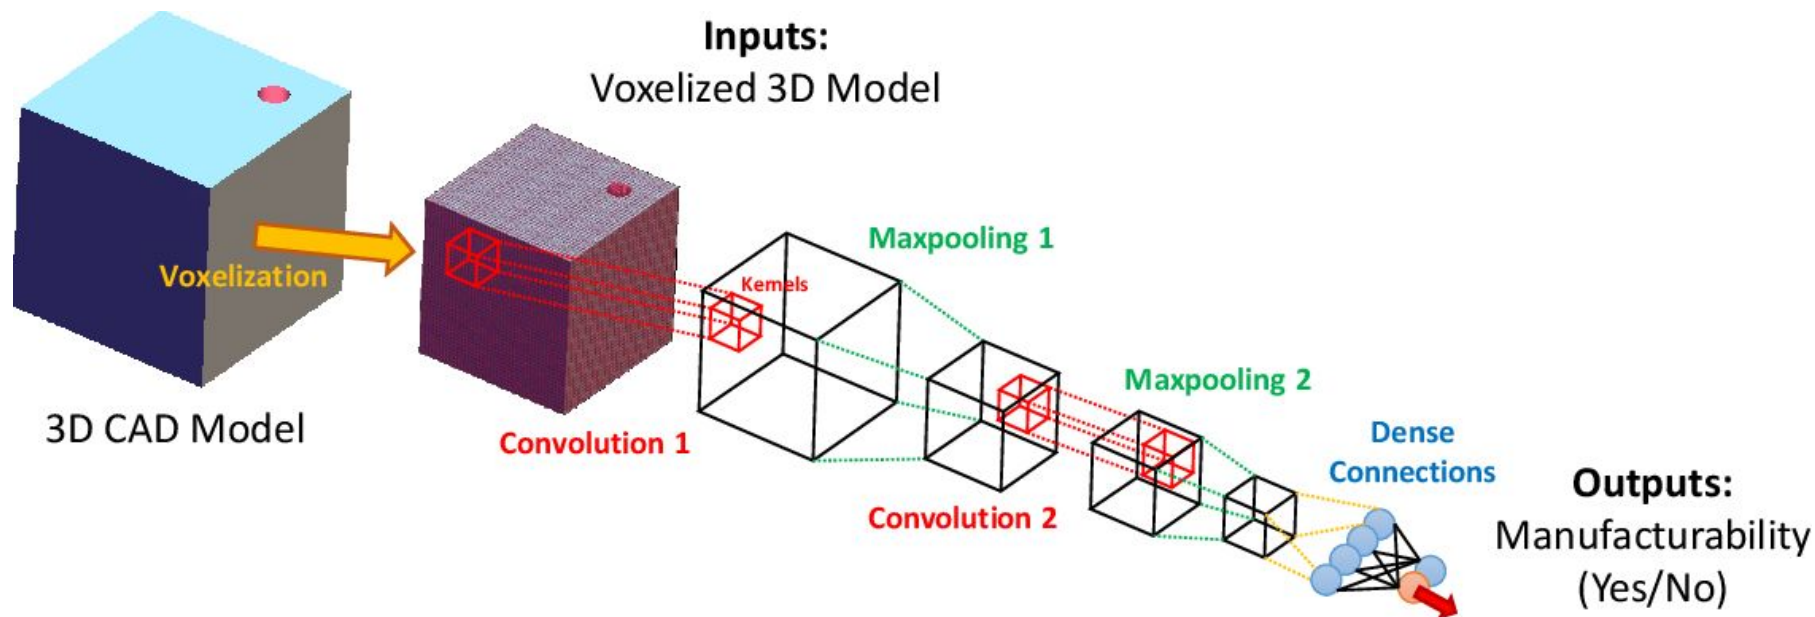


## Template Matching

- [1] Slide template over image using correlation or least square metric.
- [2] Return peak by finding min or max point of metric.
- [3] Draw bounding box using min/max point as top left corner.



# CNNs in 3D: Volumetric CNNs

Link to GitHub repository [HERE](#)



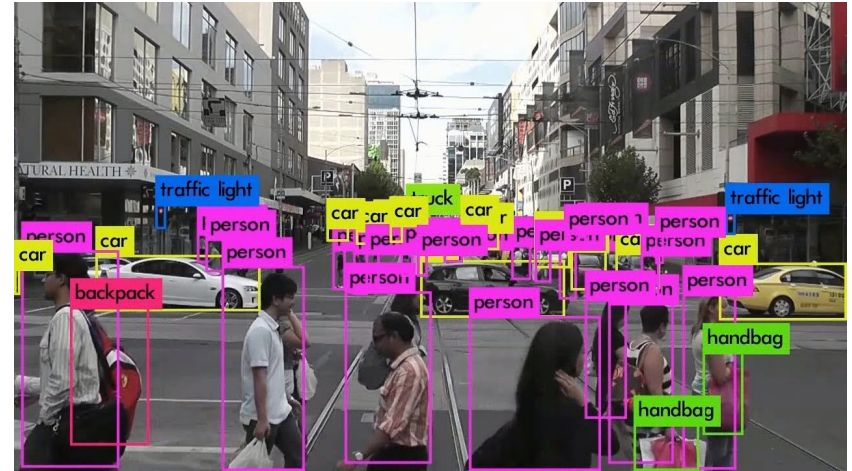
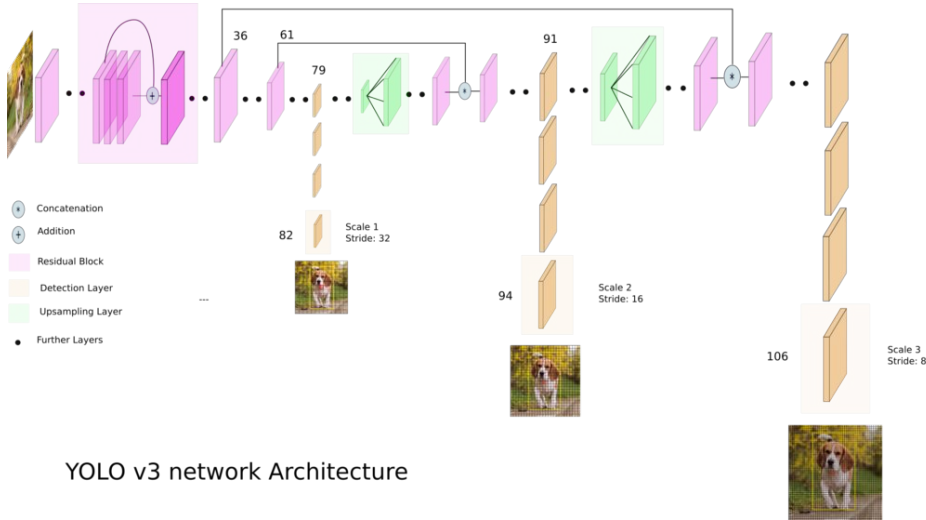
# CNNs over Time: Video Classification CNNs

Link to GitHub repository [HERE](#)



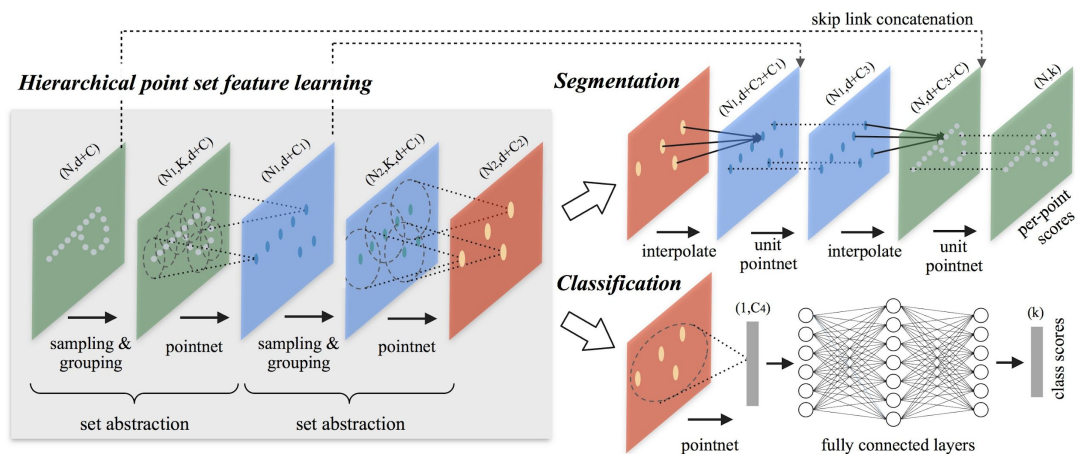
# CNNs for Self-Driving Cars: YOLO-v3

Link to GitHub repository [HERE](#)



# CNNs for Self-Driving Cars: PointNet++

Link to GitHub repository [HERE](#)



# (Some) Neural Network Architectures for Computer Vision

[ResNet](#) – Residual neural network for image classification

[VGG-16](#)/[SegNet](#) – Fully convolutional encoder-decoder networks (FCNNs)

[Inception](#) – Image classification

[AlexNet](#) – Original network for MNIST digit classification

[YOLOv3](#) – Real-time object detection and classification

[CycleGAN](#) – Image translation using GANs

[MeshCNN](#) – 3D convolutional neural networks

[PointNet++](#) – Lidar Point Cloud Object Segmentation